

Week 9 Deliverables

Group Name: Data Go

Team member's details :

Name: Yuheng Chen

Email: yuh363@gmail.com

Country: USA

College/Company: N/A

Specialization: Data Science

Name: Terry Chou

Email: techch26@gmail.com

Country: USA

College/Company: N/A

Specialization: Data Science

Name: Rishi Aluri

Email: rishialuri@gmail.com

Country: UK

College/Company: N/A

Specialization: Data Science

Name: Justine Pile

Email: justypile@gmail.com

Country: US

College/Company: N/A

Specialization: Data Science

Problem description

The objective of this project is to develop a predictive model for ABC Bank to determine the likelihood of customers subscribing to their term deposit product. By analyzing customer interactions with the bank and other financial institutions, the machine learning model will identify potential clients who are more likely to purchase the product. This will enable the bank to optimize their marketing efforts by focusing resources on customers with a higher probability of conversion, thus enhancing campaign efficiency and reducing costs. The project will assess model performance with and without using the "duration" feature while also addressing any data imbalance through suitable techniques.

Github Repo link

https://github.com/yuh39/Bank_Marketing_Group_Project/tree/main/docs/Week9

Data cleansing and transformation done on the data

Yuheng Chen – bank-full.csv:

1. Data type:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         45211 non-null  int64
1   job         45211 non-null  object
2   marital     45211 non-null  object
3   education   45211 non-null  object
4   default     45211 non-null  object
5   balance     45211 non-null  int64
6   housing     45211 non-null  object
7   loan        45211 non-null  object
8   contact     45211 non-null  object
9   day         45211 non-null  int64
10  month       45211 non-null  object
11  duration    45211 non-null  int64
12  campaign    45211 non-null  int64
13  pdays       45211 non-null  int64
14  previous    45211 non-null  int64
15  poutcome    45211 non-null  object
16  y           45211 non-null  object
dtypes: int64(7), object(10)
memory usage: 5.9+ MB
```

shape: 45211 rows × 17 columns

2. Check missing values, I replace 'unknown' values with NaN values

```
## Check missing values  
d1.isna().sum()
```

```
age          0  
job          288  
marital      0  
education    1857  
default      0  
balance      0  
housing      0  
loan         0  
contact      13020  
day          0  
month        0  
duration     0  
campaign     0  
pdays       0  
previous     0  
y            0  
dtype: int64
```

- Separate the original data set to training and testing and using random forest model to predict the missing values
- Replace only null values in the testing data from predictions and get final dataframe

```
array([[ 'management', 'tertiary', nan],  
       [ 'technician', 'secondary', nan],  
       [ 'entrepreneur', 'secondary', nan],  
       ...,  
       [ 'management', nan, 'cellular'],  
       [ 'student', nan, 'cellular'],  
       [nan, nan, 'cellular']], dtype=object)
```

⇒


```
array([[ 'management', 'tertiary', 'cellular'],  
       [ 'technician', 'secondary', 'cellular'],  
       [ 'entrepreneur', 'secondary', 'cellular'],  
       ...,  
       [ 'management', 'secondary', 'cellular'],  
       [ 'student', 'tertiary', 'cellular'],  
       [ 'management', 'secondary', 'cellular']], dtype=object)
```

- Output data with no missing values

```
df = pd.concat([d1_testing, d1_training])  
df.isna().sum()
```

```
age          0  
job          0  
marital      0  
education    0  
default      0  
balance      0  
housing      0  
loan         0  
contact      0  
day          0  
month        0  
duration     0  
campaign     0  
pdays       0  
previous     0  
y            0  
dtype: int64
```


3. We have outliers in the duration column



a. Remove outliers in duration column, the dataframe shape changed

```
Old Shape: (45211, 16)  
New Shape: (41976, 18)
```

```
: sns.boxplot(x = dl['duration'])  
: <AxesSubplot:xlabel='duration'>
```



bank-addtional.csv

Terry Chou – bank-addtional-full.csv

Data Info

#	Column	Non-Null Count	Dtype
0	age	41188 non-null	int64
1	job	41188 non-null	object
2	marital	41188 non-null	object
3	education	41188 non-null	object
4	default	41188 non-null	object
5	housing	41188 non-null	object
6	loan	41188 non-null	object
7	contact	41188 non-null	object
8	month	41188 non-null	object
9	day_of_week	41188 non-null	object
10	duration	41188 non-null	int64
11	campaign	41188 non-null	int64
12	pdays	41188 non-null	int64
13	previous	41188 non-null	int64
14	poutcome	41188 non-null	object
15	emp.var.rate	41188 non-null	float64
16	cons.price.idx	41188 non-null	float64
17	cons.conf.idx	41188 non-null	float64
18	euribor3m	41188 non-null	float64
19	nr.employed	41188 non-null	float64
20	y	41188 non-null	object

dtypes: float64(5), int64(5), object(11)

Data Problems

- Roughly **26%** of rows contains one or more “unknown” value in this dataset:

```
# rows with "unknown" proportion
len(df_bank_add_full[(df_bank_add_full.job == 'unknown') | (df_bank_add_full.marital == 'unknown') | (df_bank_add_full.education == 'unknown')])

0.25978440322424007
```

- Column “poutcome” contains **86%** of “nonexistent” value:

```
# poutcome "nonexistent" proportion
nonexistent_count = len(df_bank_add_full[(df_bank_add_full.poutcome == 'nonexistent')])
nonexistent_count/len(df_bank_add_full)

0.8634310964358551
```

- The dataset has **12 duplicate rows**:

```
duplicate = df_bank_add_full[df_bank_add_full.duplicated()]
print("Duplicate Rows :")
len(duplicate)

Duplicate Rows :

12
```

Data Cleaning Approach

- Using **Random Forest Classifier** ML model to predict the 26% “unknown” values
- Dropped the duplicate rows
- Dropped the “poutcome” column

Function Steps:

1. Replace “Unknown” with NA
2. identify columns with NAs (‘Job’, ‘marital’, ‘education’, ‘default’, ‘housing’, ‘loan’)
 - a. **df_not_missing**: rows without any NA
 - b. **df_missing**: rows with NA

```
# define independent columns and target columns
df_not_missing = df_new[~df_new[missing_columns].isna().any(axis = 1)]
df_missing = df_new[df_new[missing_columns].isnull().any(axis=1)]
```

3. Using **the rows without any NA (df_not_missing)** to train the model
 - a. **y** : target columns (‘Job’, ‘marital’, ‘education’, ‘default’, ‘housing’, ‘loan’) of **df_not_missing**
 - b. **x** : independent columns: the rest of the columns, excluding ‘y’ and ‘poutcome’, of **df_not_missing**

```
# define x variable for training data
x = df_not_missing.drop(columns = missing_columns)
x = x.drop(columns = ['y', 'poutcome'])
x = pd.get_dummies(x, columns=x.select_dtypes(include=['object']).columns)
```

```
# define y variable for training data, and the column to be filled
y = df_not_missing[column]
```

4. Apply the model on **the rows with NA (df_missing)** to predict and fill the NA values in the target columns – One column at a time.
 - a. **x_missing**: independent columns of **df_missing**
 - b. **y_missing**: target columns of **df_missing**

```
# define x variable for testing data
x_missing = df_missing.drop(columns = missing_columns)
x_missing = x_missing.drop(columns = ['y', 'poutcome'])
x_missing = pd.get_dummies(x_missing, columns=x_missing.select_dtypes(include=['
```

```
#define the column to be filled
y_missing = df_missing[column]
```

5. concatenate the **df_not_missing** rows back with the updated **df_missing** rows
6. dropped the duplicate columns
7. dropped the 'poutcome' column

Result:

Before Cleaning: Many unknown values

```
df_bank_add_full.loc[filtered_index]
```

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	...	campaign	pdays	previous	poutcome	emp.var.ra
1	57	services	married	high.school	unknown	no	no	telephone	may	mon	...	1	999	0	nonexistent	1
5	45	services	married	basic.9y	unknown	no	no	telephone	may	mon	...	1	999	0	nonexistent	1
7	41	blue-collar	married	unknown	unknown	no	no	telephone	may	mon	...	1	999	0	nonexistent	1
10	41	blue-collar	married	unknown	unknown	no	no	telephone	may	mon	...	1	999	0	nonexistent	1
15	54	retired	married	basic.9y	unknown	yes	yes	telephone	may	mon	...	1	999	0	nonexistent	1
...
41118	34	technician	married	unknown	no	yes	no	cellular	nov	tue	...	2	999	2	failure	-1
41120	60	admin.	married	unknown	no	no	no	cellular	nov	tue	...	2	999	0	nonexistent	-1
41122	34	technician	married	unknown	no	no	no	cellular	nov	tue	...	3	999	0	nonexistent	-1
41135	54	technician	married	unknown	no	yes	no	cellular	nov	thu	...	1	999	1	failure	-1
41175	34	student	single	unknown	no	yes	no	cellular	nov	thu	...	1	999	2	failure	-1

After Cleaning:

- The “unknown” value were replaced by the predicted values,
- The “poutcome” column is removed

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	duration	campaign	pdays	previous	emp.var.rate
1	57	services	married	high.school	no	no	no	telephone	may	mon	149	1	999	0	1.1
5	45	services	married	basic.9y	no	no	no	telephone	may	mon	198	1	999	0	1.1
7	41	blue-collar	married	basic.4y	no	no	no	telephone	may	mon	217	1	999	0	1.1
10	41	blue-collar	married	high.school	no	no	no	telephone	may	mon	55	1	999	0	1.1
15	54	retired	married	basic.9y	no	yes	yes	telephone	may	mon	174	1	999	0	1.1
41118	34	technician	married	university.degree	no	yes	no	cellular	nov	tue	162	2	999	2	-1.1
41120	60	admin.	married	high.school	no	no	no	cellular	nov	tue	333	2	999	0	-1.1
41122	34	technician	married	university.degree	no	no	no	cellular	nov	tue	985	3	999	0	-1.1
41135	54	technician	married	university.degree	no	yes	no	cellular	nov	thu	222	1	999	1	-1.1
41175	34	student	single	university.degree	no	yes	no	cellular	nov	thu	180	1	999	2	-1.1

No Duplicate rows

```
duplicate = df_test[df_test.duplicated()]
print("Duplicate Rows :")
len(duplicate)
```

Duplicate Rows :

0

Justine Pile - bank.csv

1. Import the csv into a dataframe

```
[2]: # Import the bank.csv file into a pandas dataframe
df = pd.read_csv("Bank_Marketing_Group_Project/dataset/bank+marketing/bank/bank.csv")
df.head()

[2]:  age;"job";"marital";"education";"default";"balance";"housing";"loan";"contact";"day";"month";"duration";"campaign";"pdays";"previous";"poutcome";"y"
0                                     30;"unemployed";"married";"primary";"no";1787;...
1                                     33;"services";"married";"secondary";"no";4789;...
2                                     35;"management";"single";"tertiary";"no";1350;...
3                                     30;"management";"married";"tertiary";"no";1476;...
4                                     59;"blue-collar";"married";"secondary";"no";0;...
```


2. Separate the data into columns using the ; as a delimiter with str.split

```
[5]: # Separate the df into separate columns using the ; as the delimiter and keep the column names as-is
df = df['age';'job';'marital';'education';'default';'balance';'housing';'loan';'contact';'day';'month';'duration';'campaign';'pdays';'previous';'poutcome';'y']
df.head()
```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
0	30	"unemployed"	"married"	"primary"	"no"	1787	"no"	"no"	"cellular"	19	"oct"	79	1	-1	0	"unknown"	"no"
1	33	"services"	"married"	"secondary"	"no"	4789	"yes"	"yes"	"cellular"	11	"may"	220	1	339	4	"failure"	"no"
2	35	"management"	"single"	"tertiary"	"no"	1350	"yes"	"no"	"cellular"	16	"apr"	185	1	330	1	"failure"	"no"
3	30	"management"	"married"	"tertiary"	"no"	1476	"yes"	"yes"	"unknown"	3	"jun"	199	4	-1	0	"unknown"	"no"
4	59	"blue-collar"	"married"	"secondary"	"no"	0	"yes"	"no"	"unknown"	5	"may"	226	1	-1	0	"unknown"	"no"

3. Rename the columns

```
[6]: #Rename the columns in df
new_column_names = [
    "age", "job", "marital", "education", "default", "balance",
    "housing", "loan", "contact", "day", "month", "duration",
    "campaign", "pdays", "previous", "poutcome", "y"
]

df.columns = new_column_names
df.head()
```

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous	poutcome	y
0	30	"unemployed"	"married"	"primary"	"no"	1787	"no"	"no"	"cellular"	19	"oct"	79	1	-1	0	"unknown"	"no"
1	33	"services"	"married"	"secondary"	"no"	4789	"yes"	"yes"	"cellular"	11	"may"	220	1	339	4	"failure"	"no"
2	35	"management"	"single"	"tertiary"	"no"	1350	"yes"	"no"	"cellular"	16	"apr"	185	1	330	1	"failure"	"no"
3	30	"management"	"married"	"tertiary"	"no"	1476	"yes"	"yes"	"unknown"	3	"jun"	199	4	-1	0	"unknown"	"no"
4	59	"blue-collar"	"married"	"secondary"	"no"	0	"yes"	"no"	"unknown"	5	"may"	226	1	-1	0	"unknown"	"no"

4. Verify there are no duplicates

```
[4]: # Verify there are no duplicate rows in dataframe
duplicate_rows = df.duplicated().sum()

if duplicate_rows == 0:
    print("No duplicate rows found in the DataFrame.")
else:
    print(f"{duplicate_rows} duplicate rows found in the DataFrame.")

No duplicate rows found in the DataFrame.
```

5. Check for null values

```
[7]: # Check for null values in df
null_values = df.isnull().sum()
print(null_values)
```

```
age          0
job          0
marital      0
education    0
default      0
balance      0
housing      0
loan         0
contact      0
day          0
month        0
duration     0
campaign     0
pdays       0
previous     0
poutcome     0
y            0
dtype: int64
```

Rishi Aluri - bank-additional.csv

1. Remove 'unkown' from dataset. Introduces NA values.

```
# Replace 'unkown' with NaN
df.replace('unknown', np.nan, inplace=True)
```

#	Column	Non-Null Count	Dtype
0	age	4119 non-null	int64
1	job	4119 non-null	object
2	marital	4119 non-null	object
3	education	4119 non-null	object
4	default	4119 non-null	object
5	housing	4119 non-null	object
6	loan	4119 non-null	object
7	contact	4119 non-null	object
8	month	4119 non-null	object
9	day_of_week	4119 non-null	object
10	duration	4119 non-null	int64
11	campaign	4119 non-null	int64
12	pdays	4119 non-null	int64
13	previous	4119 non-null	int64
14	poutcome	4119 non-null	object
15	emp.var.rate	4119 non-null	float64
16	cons.price.idx	4119 non-null	float64
17	cons.conf.idx	4119 non-null	float64
18	euribor3m	4119 non-null	float64
19	nr.employed	4119 non-null	float64
20	y	4119 non-null	object

dtypes: float64(5), int64(5), object(11)
memory usage: 675.9+ KB

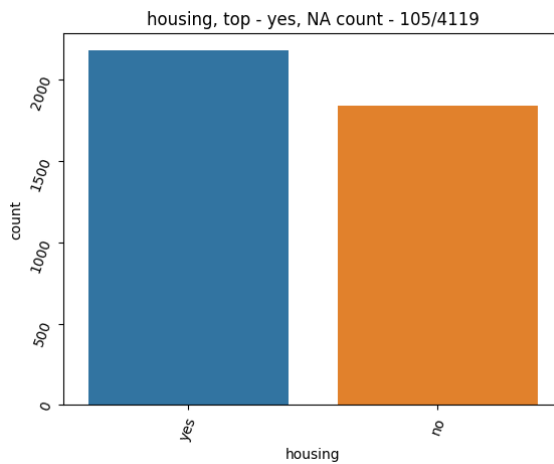
2. Cleaning categorical columns by imputing highest frequency value.

```
# Clean categorical columns - Impute Missing Values

for col in df.select_dtypes(include=['object']).columns:

    df[col].fillna(df[col].mode().values[0], inplace=True)
```

Eg-



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4119 entries, 0 to 4118
Data columns (total 21 columns):
#   Column              Non-Null Count  Dtype
---  -
0   age                  4119 non-null   float64
1   job                  4119 non-null   category
2   marital              4119 non-null   category
3   education            4119 non-null   category
4   default              4119 non-null   float64
5   housing              4119 non-null   float64
6   loan                 4119 non-null   float64
7   contact              4119 non-null   category
8   month                4119 non-null   category
9   day_of_week          4119 non-null   category
10  duration             4119 non-null   float64
11  campaign             4119 non-null   float64
12  pdays                4119 non-null   float64
13  previous             4119 non-null   float64
14  poutcome             4119 non-null   category
15  emp.var.rate         4119 non-null   float64
16  cons.price.idx       4119 non-null   float64
17  cons.conf.idx        4119 non-null   float64
18  euribor3m            4119 non-null   float64
19  nr.employed          4119 non-null   float64
20  y                    4119 non-null   float64
dtypes: category(7), float64(14)
memory usage: 480.5 KB
```

3. Column Transformations

```
# Column Transformations

for col in df.select_dtypes(include=['int64']).columns:

    df[col]=df[col].astype(np.float64)

for col in ['default', 'housing', 'loan', 'y']:

    df[col]=np.where(df[col]=='yes', 1.0, 0.0)


for col in df.select_dtypes(include=['object']).columns:

    df[col]=df[col].astype('category')
```

	count	unique	top	freq	mean	std	min	25%	50%	75%	max
age	4119.0	NaN	NaN	NaN	40.11362	10.313362	18.0	32.0	38.0	47.0	88.0
job	4119	12	admin.	1012	NaN	NaN	NaN	NaN	NaN	NaN	NaN
marital	4119	4	married	2509	NaN	NaN	NaN	NaN	NaN	NaN	NaN
education	4119	8	university.degree	1264	NaN	NaN	NaN	NaN	NaN	NaN	NaN
default	4119	3	no	3315	NaN	NaN	NaN	NaN	NaN	NaN	NaN
housing	4119	3	yes	2175	NaN	NaN	NaN	NaN	NaN	NaN	NaN
loan	4119	3	no	3349	NaN	NaN	NaN	NaN	NaN	NaN	NaN
contact	4119	2	cellular	2652	NaN	NaN	NaN	NaN	NaN	NaN	NaN
month	4119	10	may	1378	NaN	NaN	NaN	NaN	NaN	NaN	NaN
day_of_week	4119	5	thu	860	NaN	NaN	NaN	NaN	NaN	NaN	NaN
duration	4119.0	NaN	NaN	NaN	256.788055	254.703736	0.0	103.0	181.0	317.0	3643.0
campaign	4119.0	NaN	NaN	NaN	2.537266	2.568159	1.0	1.0	2.0	3.0	35.0
pdays	4119.0	NaN	NaN	NaN	960.42219	191.922786	0.0	999.0	999.0	999.0	999.0
previous	4119.0	NaN	NaN	NaN	0.190337	0.541788	0.0	0.0	0.0	0.0	6.0
poutcome	4119	3	nonexistent	3523	NaN	NaN	NaN	NaN	NaN	NaN	NaN
emp.var.rate	4119.0	NaN	NaN	NaN	0.084972	1.563114	-3.4	-1.8	1.1	1.4	1.4
cons.price.idx	4119.0	NaN	NaN	NaN	93.579704	0.579349	92.201	93.075	93.749	93.994	94.767
cons.conf.idx	4119.0	NaN	NaN	NaN	-40.499102	4.594578	-50.8	-42.7	-41.8	-36.4	-26.9
euribor3m	4119.0	NaN	NaN	NaN	3.621356	1.733591	0.635	1.334	4.857	4.961	5.045
nr.employed	4119.0	NaN	NaN	NaN	5166.481695	73.667904	4963.6	5099.1	5191.0	5228.1	5228.1
y	4119	2	no	3668	NaN	NaN	NaN	NaN	NaN	NaN	NaN

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4119 entries, 0 to 4118
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   age                   4119 non-null   float64
1   job                   4119 non-null   category
2   marital               4119 non-null   category
3   education             4119 non-null   category
4   default               4119 non-null   float64
5   housing               4119 non-null   float64
6   loan                  4119 non-null   float64
7   contact               4119 non-null   category
8   month                 4119 non-null   category
9   day_of_week           4119 non-null   category
10  duration               4119 non-null   float64
11  campaign               4119 non-null   float64
12  pdays                 4119 non-null   float64
13  previous               4119 non-null   float64
14  poutcome              4119 non-null   category
15  emp.var.rate          4119 non-null   float64
16  cons.price.idx         4119 non-null   float64
17  cons.conf.idx          4119 non-null   float64
18  euribor3m             4119 non-null   float64
19  nr.employed           4119 non-null   float64
20  y                     4119 non-null   float64
dtypes: category(7), float64(14)
memory usage: 480.5 KB
```

eg-



4. Cleaning numerical values by normalizing:

```
# Clean numerical columns - Normalize using MinMax
for col in df.select_dtypes(include=['float64']).columns:
    min_value = df[col].min()
    max_value = df[col].max()
    df[col] = ((df[col] - min_value) / (max_value - min_value))
```

eg-

